



Credit Card Fraud Detection

Under Supervision of: Mr. Ubaid Zaman

INTRODUCTION TO DATA SCIENCE

AMIN M QURAISHI (63855) & BUSHRA MUNEER (63759)

Contents

Importing Libraries

Adding files and loading dataset into data frame

First five data of dataset

Last five data of dataset

Information of dataset

Check for the missing values in each column

Distribution of normal and fraud transactions

Separating the data

Comparing the values of both normal and fraud transactions

Under sampling

K-Neighbors Classifier

Logistic Regression

Decision Tree

Bagging Classifier (base as Logistic Regression)

Random Forest Classifier

Summary

▼ Credit Card Fraud Detection

Importing Libraries

```
[ ] from google.colab import drive
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, recall_score, precision_score
from sklearn.metrics import confusion_matrix, classification_report
from sklearn.metrics import plot_confusion_matrix
from sklearn.tree import DecisionTreeClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn import model_selection
from sklearn import tree
from sklearn.neighbors import KNeighborsClassifier
```

Adding Files and loading dataset into dataframe

```
[ ] drive.mount('/content/drive')
df = pd.read_csv('/content/drive/My Drive/creditcard.csv')
```

Mounted at /content/drive

First five data of dataset

```
[ ] df.head()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...

5 rows × 31 columns



last five data of dataset

df.tail()

	Time	V1	V2	V3	V4	V5	V6	V7	V8	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.9
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.5
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.4
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.3
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.4

5 rows x 31 columns

Information of dataset

```
[ ] df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    Time    284807 non-null  float64
1    V1       284807 non-null  float64
2    V2       284807 non-null  float64
3    V3       284807 non-null  float64
4    V4       284807 non-null  float64
5    V5       284807 non-null  float64
6    V6       284807 non-null  float64
7    V7       284807 non-null  float64
8    V8       284807 non-null  float64
9    V9       284807 non-null  float64
10   V10      284807 non-null  float64
11   V11      284807 non-null  float64
12   V12      284807 non-null  float64
13   V13      284807 non-null  float64
14   V14      284807 non-null  float64
15   V15      284807 non-null  float64
16   V16      284807 non-null  float64
17   V17      284807 non-null  float64
18   V18      284807 non-null  float64
19   V19      284807 non-null  float64
20   V20      284807 non-null  float64
21   V21      284807 non-null  float64
22   V22      284807 non-null  float64
23   V23      284807 non-null  float64
24   V24      284807 non-null  float64
25   V25      284807 non-null  float64
26   V26      284807 non-null  float64
27   V27      284807 non-null  float64
28   V28      284807 non-null  float64
29   Amount   284807 non-null  float64
30   Class    284807 non-null  int64   
dtypes: float64(30), int64(1)
memory usage: 67.4 MB
```

Check for the missing values in each column

```
[ ] df.isnull().sum()
```

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       0
V20       0
V21       0
V22       0
V23       0
V24       0
V25       0
V26       0
V27       0
V28       0
Amount    0
Class     0
dtype: int64
```

Distribution of normal and fraud transactions

```
[ ] df['Class'].value_counts()
```

```
0    284315
1      492
Name: Class, dtype: int64
```

dataset highly unbalanced..

0 normal transaction, 1 fraudulent transaction

Separating the data

```
[ ] non_fraud=df[df.Class == 0]
    fraud=df[df.Class == 1]
```

```
[ ] print(non_fraud.shape)
    print(fraud.shape)
```

```
(284315, 31)
(492, 31)
```

Comparing the values of both normal and fraud transactions

```
[ ] df.groupby('Class').mean()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8
Class									
0	94838.202258	0.008258	-0.006271	0.012171	-0.007860	0.005453	0.002419	0.009637	-0.000987
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.568731	0.570636

2 rows x 30 columns

Under sampling

Building a sample dataset which will be containing similar distribution of normal transactions and Fraudulent Transactions

Number of Fraudulent Transactions : 492


```
[ ] non_fraud_sample = non_fraud.sample(n=492)
```

```
[ ] new_dataset = pd.concat([non_fraud_sample,fraud],axis=0)
```


```
[ ] new_dataset.head()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8
102920	68427.0	-0.312981	0.505159	1.603193	0.701377	-0.195373	-0.521946	0.280061	0.038075
84973	60544.0	1.103947	0.219566	0.155744	1.027188	-0.051178	-0.400656	0.212803	0.000572
35168	37998.0	1.046539	-0.407327	0.919160	0.649235	-1.184232	-0.663533	-0.459903	0.030745
136323	81648.0	-0.173644	-2.953591	-0.672841	0.080156	-1.715645	-0.826130	0.830937	-0.552689
242982	151741.0	-0.055903	-0.976752	0.030459	-0.863338	0.548546	-0.533087	-0.619720	0.048248

5 rows x 31 columns

 new_dataset.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 984 entries, 102920 to 281674
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype  
---  -
0    Time        984 non-null    float64
1    V1          984 non-null    float64
2    V2          984 non-null    float64
3    V3          984 non-null    float64
4    V4          984 non-null    float64
5    V5          984 non-null    float64
6    V6          984 non-null    float64
7    V7          984 non-null    float64
8    V8          984 non-null    float64
9    V9          984 non-null    float64
10   V10         984 non-null    float64
11   V11         984 non-null    float64
12   V12         984 non-null    float64
13   V13         984 non-null    float64
14   V14         984 non-null    float64
15   V15         984 non-null    float64
16   V16         984 non-null    float64
17   V17         984 non-null    float64
18   V18         984 non-null    float64
19   V19         984 non-null    float64
20   V20         984 non-null    float64
21   V21         984 non-null    float64
22   V22         984 non-null    float64
23   V23         984 non-null    float64
24   V24         984 non-null    float64
25   V25         984 non-null    float64
26   V26         984 non-null    float64
27   V27         984 non-null    float64
28   V28         984 non-null    float64
29   Amount      984 non-null    float64
30   Class       984 non-null    int64  
dtypes: float64(30), int64(1)
memory usage: 246.0 KB
```

 new_dataset.tail()

```
<img alt="copy icon" data-bbox="138 672 158 685"/>
      Time      V1      V2      V3      V4      V5      V6      V7      V8
279863  169142.0 -1.927883  1.125653 -4.518331  1.749293 -1.566487 -2.010494 -0.882850  0.697211 -2.064
280143  169347.0  1.378559  1.289381 -5.004247  1.411850  0.442581 -1.326536 -1.413170  0.248525 -1.127
280149  169351.0 -0.676143  1.126366 -2.213700  0.468308 -1.120541 -0.003346 -2.234739  1.210158 -0.652
281144  169966.0 -3.113832  0.585864 -5.399730  1.817092 -0.840618 -2.943548 -2.208002  1.058733 -1.632
281674  170348.0  1.991976  0.158476 -2.583441  0.408670  1.151147 -0.096695  0.223050 -0.068384  0.577

5 rows x 31 columns
```



```
new_dataset['Class'].value_counts()
```

```
0    492
1    492
Name: Class, dtype: int64
```

```
X=new_dataset.drop(columns='Class',axis=1)
Y=new_dataset['Class']
```

```
print(X)
```

```
Time      V1      V2      V3      V4      V5      V6  \
102920  68427.0 -0.312981  0.505159  1.603193  0.701377 -0.195373 -0.521946
84973   60544.0  1.103947  0.219566  0.155744  1.027188 -0.051178 -0.400656
35168   37998.0  1.046539 -0.407327  0.919160  0.649235 -1.184232 -0.663533
136323  81648.0 -0.173644 -2.953591 -0.672841  0.080156 -1.715645 -0.826130
242982  151741.0 -0.055903 -0.976752  0.030459 -0.863338  0.548546 -0.533087
...      ...      ...      ...      ...      ...      ...
279863  169142.0 -1.927883  1.125653 -4.518331  1.749293 -1.566487 -2.010494
280143  169347.0  1.378559  1.289381 -5.004247  1.411850  0.442581 -1.326536
280149  169351.0 -0.676143  1.126366 -2.213700  0.468308 -1.120541 -0.003346
281144  169966.0 -3.113832  0.585864 -5.399730  1.817092 -0.840618 -2.943548
281674  170348.0  1.991976  0.158476 -2.583441  0.408670  1.151147 -0.096695

V7      V8      V9      ...      V20      V21      V22  \
102920  0.280061  0.038075  0.156139  ... -0.029140  0.037971  0.277495
84973   0.212803  0.000572 -0.352222  ... -0.149483  0.074633  0.177438
35168   -0.459903  0.030745  0.709154  ...  0.010701  0.033611 -0.160247
136323  0.830937 -0.552689 -1.170958  ...  1.246760  0.118773 -0.984116
242982 -0.619720  0.048248 -0.547072  ...  0.473728  0.510498  1.355643
...      ...      ...      ...      ...      ...      ...
279863 -0.882850  0.697211 -2.064945  ...  1.252967  0.778584 -0.319189
280143 -1.413170  0.248525 -1.127396  ...  0.226138  0.370612  0.028234
280149 -2.234739  1.210158 -0.652250  ...  0.247968  0.751826  0.834108
281144 -2.208002  1.058733 -1.632333  ...  0.306271  0.583276 -0.269209
281674  0.223050 -0.068384  0.577829  ... -0.017652 -0.164350 -0.295135

V23      V24      V25      V26      V27      V28  Amount
102920 -0.119164  0.433383 -0.391840  0.435189  0.132972  0.154542    4.99
84973  -0.070577  0.210336  0.574347 -0.359584  0.008746  0.008916   30.76
35168  0.001088  0.386509  0.068980  0.386655 -0.033160  0.039037   95.70
136323 -0.686252  0.495051  0.006298  0.949754 -0.214172  0.149118  825.73
242982  0.349074  0.570515 -1.661312 -0.240226  0.148534  0.131473   55.00
...      ...      ...      ...      ...      ...      ...
279863  0.639419 -0.294885  0.537503  0.788395  0.292680  0.147968  390.00
280143 -0.145640 -0.081049  0.521875  0.739467  0.389152  0.186637    0.76
280149  0.190944  0.032070 -0.739695  0.471111  0.385107  0.194361   77.89
281144 -0.456108 -0.183659 -0.328168  0.606116  0.884876 -0.253700  245.00
281674 -0.072173 -0.450261  0.313267 -0.289617  0.002988 -0.015309   42.53
```

```
[984 rows x 30 columns]
```

```
print(Y)
```

```
102920    0
84973     0
35168     0
136323    0
242982    0
...
279863    1
280143    1
280149    1
281144    1
281674    1
Name: Class, Length: 984, dtype: int64
```


Splitting data to test and train

```
[ ] X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.25, stratify=Y, random_state=1)

[ ] print(X.shape, X_train.shape, X_test.shape)

(984, 30) (738, 30) (246, 30)
```

KNeighbors Classifier


K-NN is a Supervised Learning Technique which assumes the similarity b/w new and available data and put the new data into the category that is most similar to the available categories.

```
[ ] classifier = KNeighborsClassifier()
classifier.fit(X_train, Y_train)
X_test_prediction = classifier.predict(X_test)
X_train_prediction = classifier.predict(X_train)
```

Train Data

```
[ ] print(confusion_matrix(Y_train, X_train_prediction))
plot_confusion_matrix(classifier,X_train,Y_train)
plt.show()
print(classification_report(Y_train, X_train_prediction))
print('Accuracy Using KNeighbors classifier on Train Data is ',accuracy_score(X_train_prediction,Y_train)*100)
```

```
[[288  81]
 [ 82 287]]
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated.
warnings.warn(msg, category=FutureWarning)
```




Accuracy Using KNeighbors classifier on Train Data is 77.91327913279133

```
[ ] print(confusion_matrix( Y_test,X_test_prediction))
plot_confusion_matrix(classifier,X_test,Y_test)
plt.show()
print(classification_report(Y_test, X_test_prediction))
```

```
[[68 55]
 [53 70]]
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated.
warnings.warn(msg, category=FutureWarning)
```



```
▶ print('KNeighbors classifier Accuracy -> ',accuracy_score(X_test_prediction,Y_test)*100)
print("KNeighbors classifier Precision Score -> ",precision_score(X_test_prediction, Y_test,average = 'weighted')*100)
print("KNeighbors classifier Recall Score -> ",recall_score(X_test_prediction, Y_test,average = 'weighted')*100)
```

```
☞ KNeighbors classifier Accuracy -> 60.16260162601627
KNeighbors classifier Precision Score -> 60.16260162601627
KNeighbors classifier Recall Score -> 60.16260162601627
```


↳ Logistic Regression

Logistic Regression is a Supervised Learning Technique and is used for solving the Classification problems, predict the categorical dependent variable with the help of independent variables.

```
[ ] classifier = LogisticRegression()  
    classifier.fit(X_train, Y_train)
```

```
▶ X_train_prediction = classifier.predict(X_train)  
print(confusion_matrix( Y_train,X_train_prediction))  
plot_confusion_matrix(classifier,X_train,Y_train)  
plt.show()  
print(classification_report( X_train_prediction,Y_train))  
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

```
↳ [[353  16]  
   [ 28 341]]  
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated.  
warnings.warn(msg, category=FutureWarning)
```



```
[ ] print('Accuracy Using Logistic Regression on Training data : ', (training_data_accuracy*100))
```

```
↳ Accuracy Using Logistic Regression on Training data : 94.03794037940379
```

```

X_test_prediction = classifier.predict(X_test)
print(confusion_matrix(Y_test,X_test_prediction))
plot_confusion_matrix(classifier,X_test,Y_test)
plt.show()
print(classification_report(Y_test, X_test_prediction))
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)


```

```

[[112  11]
 [  9 114]]

```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated.
warnings.warn(msg, category=FutureWarning)



```

print('Logistic Regression Accuracy Score -> ', test_data_accuracy*100)
print("Logistic Regression Precision Score -> ",precision_score(X_test_prediction, Y_test,average = 'weighted')*100)
print("Logistic Regression Recall Score -> ",recall_score(X_test_prediction, Y_test,average = 'weighted')*100)

```

```

Logistic Regression Accuracy Score -> : 91.869918699187
Logistic Regression Precision Score -> 60.16260162601627
Logistic Regression Recall Score -> 60.16260162601627

```

▼ Decision Tree

Decision Tree is a Supervised learning technique based on a tree-structured classifier, where internal nodes represent the features of a dataset, branches represent the decision rules and each leaf node represents the outcome. There are two nodes, which are the Decision Node and Leaf Node. Decision nodes are used to make any decision and have multiple branches, whereas Leaf nodes are the output of those decisions and do not contain any further branches. The decisions or the test are performed on the basis of features of the given dataset.


```
accuracy=[]
recall_scores=[]
max_depths=[2,4,6,8,10,12,14]
for n in max_depths:
    clf_u=DecisionTreeClassifier(max_depth=n)
    clf_u.fit(X_train,Y_train)
    predict=clf_u.predict(X_test)
    recall_scores.append(recall_score(Y_test,predict))
    accuracy.append(clf_u.score(X_test,Y_test))
```

```
[ ] fig,(ax1,ax2)=plt.subplots(1,2,figsize=(12,5))
```

```
figs=[ax1,ax2]
```

```
def plot_graphs(fig_index,y,y_label,X=max_depths):
    global figs
    figs[fig_index].plot(X,y)
    figs[fig_index].set_ylabel(y_label)
print(accuracy)
plot_graphs(1,accuracy,'accuracy')
plot_graphs(0,recall_scores,'recall')
```

```
[0.9105691056910569, 0.8983739837398373, 0.9024390243902439, 0.8658536585365854, 0.8617886178861789, 0.8699186991869918, 0.8
```




```

[ ] max_depth=max_depths[recall_scores.index(max(recall_scores))]
    clf_u=DecisionTreeClassifier(max_depth=max_depth)
    clf_u.fit(X_train,Y_train)
    predict=clf_u.predict(X_test)
    recall_score(Y_test,predict)
    print(confusion_matrix( Y_test,predict))
    plot_confusion_matrix(clf_u,X_test,Y_test)
    plt.show()
    print(classification_report(Y_test, predict))
    test_data_accuracy = accuracy_score(predict, Y_test)
    print('Accuracy Using Decision Tree on Test Data : ', test_data_accuracy*100)
    print('Max_depth with maximum recall {}'.format(max_depth))


```

```

[[115  8]
 [ 14 109]]

```

/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated.
warnings.warn(msg, category=FutureWarning)



	precision	recall	f1-score	support
0	0.89	0.93	0.91	123
1	0.93	0.89	0.91	123
accuracy			0.91	246
macro avg	0.91	0.91	0.91	246
weighted avg	0.91	0.91	0.91	246

Accuracy Using Decision Tree on Test Data : 91.05691056910568
Max_depth with maximum recall 2

```

▶ # Use accuracy_score function to get the accuracy
print('Decision Tree Accuracy Score : ', test_data_accuracy*100)
print("Decision Tree Precision Score -> ",precision_score(predict, Y_test,average = 'weighted')*100)
print("Decision Tree Recall Score -> ",recall_score(predict, Y_test,average = 'weighted')*100)

```

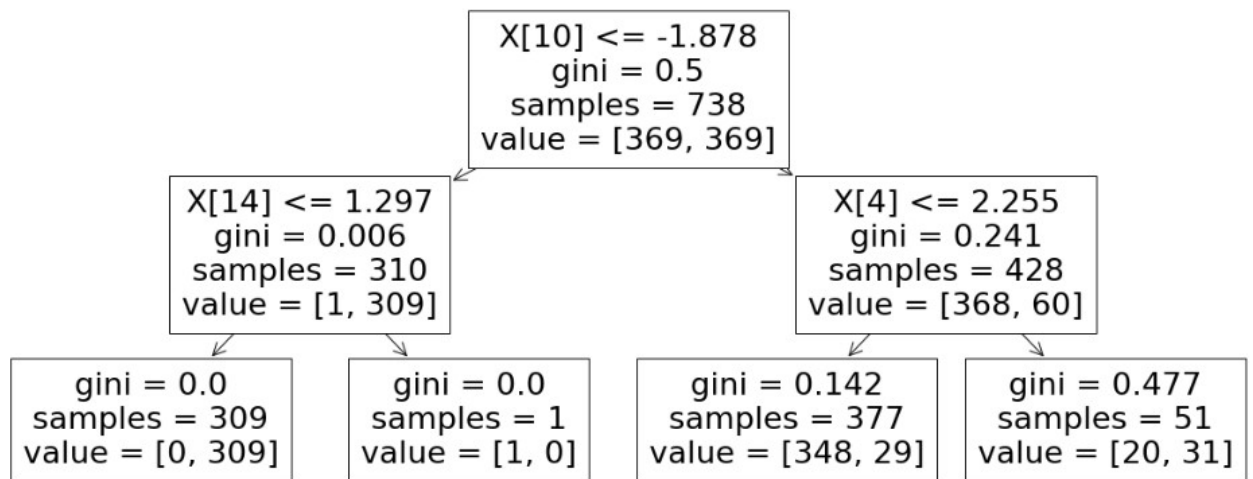
```

❏ Decision Tree Accuracy Score : 91.05691056910568
Decision Tree Precision Score -> 91.17588736862979
Decision Tree Recall Score -> 91.05691056910568

```

```
[ ] import sklearn  
plt.figure(figsize=(20,8))  
sklearn.tree.plot_tree(clf_u,max_depth=3);
```

↳




▼ Bagging Classifier (base as Logistic Regression)

A Bagging classifier is an ensemble meta-estimator that fits base classifiers each on random subsets of the original dataset and then aggregate their individual predictions (either by voting or by averaging) to form a final prediction

```
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import make_pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import BaggingClassifier
pipeline = make_pipeline(StandardScaler(), LogisticRegression())
bgclassifier = BaggingClassifier(base_estimator=pipeline)
predictions = bgclassifier.fit(X_train, Y_train)
predictions = bgclassifier.predict(X_test)
print(confusion_matrix(predictions, Y_test))
plot_confusion_matrix(bgclassifier, X_test, Y_test)
plt.show()
print('Model test Score: %.3f, ' % bgclassifier.score(X_test, Y_test),
      'Model training Score: %.3f' % bgclassifier.score(X_train, Y_train))
```

```
[ ] [[119 11]
     [ 4 112]]
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated
warnings.warn(msg, category=FutureWarning)
```



Model test Score: 0.939, Model training Score: 0.950

```
[ ] print('Accuracy Using Bagging Classifier with Logistic Regression on Training data : ', (bgclassifier.score(X_train, Y_train)*100))
print('Accuracy Using Bagging Classifier with Logistic Regression on Test Data : ', bgclassifier.score(X_test, Y_test)*100)
```

```
[ ] Accuracy Using Bagging Classifier with Logistic Regression on Training data : 94.98644986449864
Accuracy Using Bagging Classifier with Logistic Regression on Test Data : 93.90243902439023
```

```
[ ] print("Bagging Classifier with Logistic Regression Accuracy : ", accuracy_score(predictions, Y_test)*100)
print("Bagging Classifier with Logistic Regression Precision : ", precision_score(predictions, Y_test, average = 'weighted')*100)
print("Bagging Classifier with Logistic Regression Recall : ", recall_score(predictions, Y_test, average = 'weighted')*100)
```


```
Bagging Classifier with Logistic Regression Accuracy : 93.90243902439023
Bagging Classifier with Logistic Regression Precision : 94.06437966818693
Bagging Classifier with Logistic Regression Recall : 93.90243902439023
```

Random Forest Classifier

The Random forest classifier creates a set of decision trees from a randomly selected subset of the training set. It is basically a set of decision trees (DT) from a randomly selected subset of the training set and then it collects the votes from different decision trees to decide the final prediction.

```
from sklearn import metrics
classifier = RandomForestClassifier(n_estimators = 100)
classifier.fit(X_train, Y_train)
y_pred = classifier.predict(X_test)
recall_score(Y_test,y_pred)
print(confusion_matrix(y_pred,Y_test))
plot_confusion_matrix(classifier,X_test,Y_test)
plt.show()
print(classification_report(Y_test, y_pred))
print('Accuracy Using Random Forest Classifier', metrics.accuracy_score(Y_test, y_pred)*100)
```

```
[[121 13]
 [ 2 110]]
/usr/local/lib/python3.7/dist-packages/sklearn/utils/deprecation.py:87: FutureWarning: Function plot_confusion_matrix is deprecated; Function `plot_confus:
warnings.warn(msg, category=FutureWarning)
```



	precision	recall	f1-score	support
0	0.90	0.98	0.94	123
1	0.98	0.89	0.94	123
accuracy			0.94	246
macro avg	0.94	0.94	0.94	246
weighted avg	0.94	0.94	0.94	246

Accuracy Using Random Forest Classifier 93.90243902439023

```
print("Random forest classifier Accuracy : ", accuracy_score(y_pred,Y_test)*100)
print("Random forest classifier Precision : ", precision_score(y_pred,Y_test,average = 'weighted')*100)
print("Random forest classifier Recall : ", recall_score(y_pred,Y_test,average = 'weighted')*100)
```

```
Random forest classifier Accuracy : 93.90243902439023
Random forest classifier Precision : 94.3023332672351
Random forest classifier Recall : 93.90243902439023
```

▼ Summary

Comparison Of Algorithms

S.No	Algorithms	Accuracy in %	Precision in %	Recall in %
1	K Nearest Neighbors	60	60	60
2	Logistic Regression	92	60	60
3	Decision Tree	91	91	91
4	Bagging Classifier	93	94	93
5	Random Forest Classifier	93	94	93

So in order to Conclude, It is identified that bagging classifier and Random forest performs well as it was able to tell that how many times the ML model was correct overall, how good the model is at predicting a specific category and how many times the model was able to detect a specific category where as in Logistic regression the model wasnot performed very well at predicting and detecting number of times the specific category.

Random Forest Classifier/Bagging Classifier > Decision Tree > Logistic Regression > K Nearest Neighbor